

# Fraudulent Firm Classification: A Case Study of an External Audit

## Project description:

Please read the Data Set Information section to learn about this dataset. Data description is also provided for this dataset. Read data into Jupyter notebook, use pandas to import data into a data frame. Preprocess data: Explore data, check for missing data and apply data scaling. Justify the type of scaling used.

## Objectives of the project

1) Used Linear Regression, Polynomial Regression and KNN Regressor to determine the Risk Audit Score for 777 target firms. 2) Built a classification model to predict the Risk Audit Class (Fraud or Not-Fraud) of the firms with 97% accuracy on test score. 3) Trained several machine learning algorithms for classification (Decision Tree, KNN, Logistic Regression and SVC) to find the best model based on accuracy and performance.

## Regression Task:

We used several regression models to find the predict the audit score on the test set and used Grid Search to find the best scaling parameter. Also used plots and graphs to get a better glimpse of the results. Finally we reported the best regressor for our dataset.

## Classification task:

We used several classification models: KNN classification, Logistic Regression, Linear Support Vector Machine, Kernelized Support Vector Machine, Decision Tree, to classify whether a company is indulging in fraudulent practices or not and reported the best model based on scores.

## Data Set Information:

This dataset is taken from a research explained here.

The goal of the research is to help the auditors by building a classification model that can predict the fraudulent firm on the basis of the present and historical risk factors. The information about the sectors and the counts of firms are listed respectively as Irrigation (114), Public Health (77), Buildings and Roads (82), Forest (70), Corporate (47), Animal Husbandry (95), Communication (1), Electrical (4), Land (5), Science and Technology (3), Tourism (1), Fisheries (41), Industries (37), Agriculture (200).

There are two csv files to present data. Please merge these two datasets into one dataframe.

***The main objective of this project is to perform the audit risk analysis using 776 target firm's historical data. Our main focus is on determining the Risk Audit Score and prediction of Risk Class. We used several supervised techniques (Regression, Classification) to determine the Risk Audit Score and Predict the Risk Class which will be discussed below.***

***We followed a sequence of steps starting with importing required libraries, data merging, data pre-processing, data visualization, etc. So let's get started***

## Importing Libraries

```
In [15]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
import seaborn as sns
```

## Importing Datasets

```
In [16]: audit = pd.read_csv('audit_risk.csv')
trial = pd.read_csv('trial.csv')
```


## Data Distribution

```
In [17]: sns.countplot(audit['Risk'])
audit.groupby('Risk').count()
```

Out[17]:

	Sector_score	LOCATION_ID	PARA_A	Score_A	Risk_A	PARA_B	Score_B	Risk_B	TOTAL
Risk									
0	471	471	471	471	471	471	471	471	471
1	305	305	305	305	305	305	305	305	305

2 rows × 26 columns




### Audit\_Risk Vs Risk

```
In [18]: plt.scatter(audit['Audit_Risk'], audit['Risk'])


## We can see an outlier that screws can screw out analysis. We will handle this outlier in data pre-processing.
```

Out[18]: <matplotlib.collections.PathCollection at 0x1e46a13cd68>




```
In [19]: sns.boxplot(audit['numbers'],audit['TOTAL'])
```

```
Out[19]: <matplotlib.axes._subplots.AxesSubplot at 0x1e46a0c5898>
```



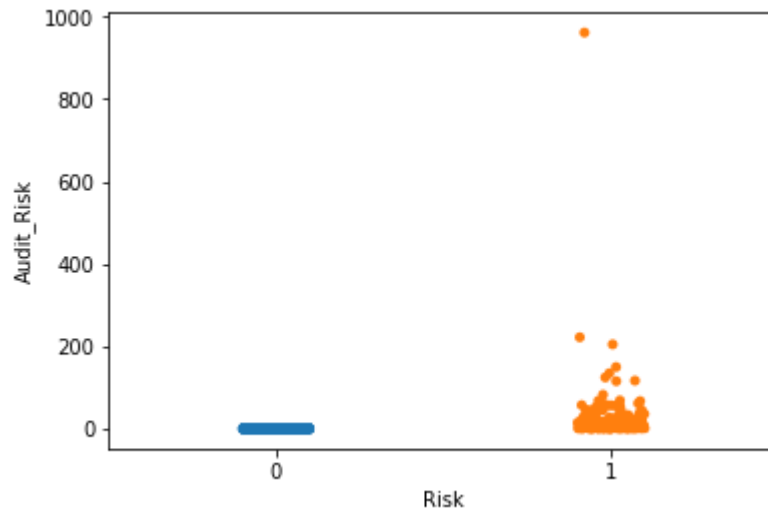
```
In [20]: plt.figure(figsize=(30,20))
sns.barplot(x='LOCATION_ID',y='Audit_Risk',data= audit,estimator=np.mean)
```

```
Out[20]: <matplotlib.axes._subplots.AxesSubplot at 0x1e46a16cb70>
```



```
In [21]: #Looking for outliers in the audit's risk column  
sns.stripplot(x='Risk',y='Audit_Risk',data=audit)
```

```
Out[21]: <matplotlib.axes._subplots.AxesSubplot at 0x1e46a383c88>
```



## Data Pre-processing

1. Imputing Missing values
2. Merging Data
3. Checking for outliers

### 1. Imputing Missing Values

```
In [22]: audit['Money_Value'].unique()
```

```
Out[22]: array([3.3800e+00, 9.4000e-01, 0.0000e+00, 1.1750e+01, 2.9500e+00,
4.4950e+01, 7.7900e+00, 7.3400e+00, 1.9300e+00, 4.4200e+00,
9.6000e-01, 1.0430e+01, 7.0000e-03, 9.0000e+00, 4.1280e+01,
1.4030e+01, 6.3180e+01, 3.4240e+01, 1.0000e-02, 2.0519e+02,
1.0000e-01, 1.1160e+01, 1.2500e+00, 1.4600e+00, 6.7800e+00,
1.1600e+00, 1.5241e+02, 1.0800e+00, 2.8400e+00, 9.0000e-01,
9.6700e+00, 3.2680e+01, 9.3503e+02, 2.9630e+01, 1.1000e-01,
2.6200e+00, 6.0000e-02, 2.4300e+00, 1.2613e+02, 2.0790e+01,
1.5692e+02, 1.2290e+01, 2.2900e+00, 7.7800e+00, 2.5100e+00,
8.3100e+00, 1.6000e-01, 4.7900e+01, 8.9100e+00, 4.9500e+00,
1.7500e+00, 6.8000e-01, 1.5820e+01, 5.8000e-01, 2.1531e+02,
5.0000e-02, 5.3340e+01, 1.0690e+01, 5.6900e+00, 1.3500e+00,
1.1690e+01, 1.4600e+01, 2.0780e+01, 1.0222e+02, 7.4000e-01,
8.7337e+02, 1.4000e+00, 4.2000e-01, 3.6520e+01, 6.0200e+00,
1.7160e+01, 1.2910e+01, 1.0790e+01, 3.4600e+00, 2.3300e+00,
5.5800e+00, 7.6000e-01, 2.5270e+01, 2.0000e-01, 7.5600e+00,
6.7030e+01, 1.9400e+01, 3.4830e+01, 9.8750e+01, 3.6000e-01,
3.5210e+01, 8.2000e+00, 2.0330e+01, 3.5130e+01, 6.3700e+01,
2.8000e-01, 2.7130e+01, 1.0270e+01, 1.3050e+01, 1.8790e+01,
1.6820e+01, 1.0030e+01, 1.2670e+01, 1.3790e+01, 1.0160e+01,
9.7300e+00, 2.6950e+01, 2.9070e+01, 9.7600e+00, 1.3310e+01,
7.7000e+00, 1.0400e+01, 2.0620e+01, 4.4020e+01, 4.3530e+01,
1.2030e+01, 1.1880e+01, 6.9600e+00, 2.1070e+01, 5.7160e+01,
1.0650e+01, 5.3600e+00, 2.4500e+00, 3.1610e+01, 3.4320e+01,
4.2400e+00, 8.3800e+00, 2.2100e+01, 5.4070e+01, 2.7680e+01,
2.1450e+01, 1.1090e+01, 1.0140e+01, 5.2130e+01, 1.7020e+01,
1.8000e+01, 6.9760e+01, 1.3880e+01, 4.0300e+00, 2.4400e+00,
7.8980e+01, 6.8800e+00, 5.4900e+00, 3.7790e+01, 6.0880e+01,
1.8320e+02, 1.0020e+01, 3.2000e-01, 1.8000e-01, 4.8560e+01,
1.8450e+01, 8.4400e+00, 3.9650e+01, 9.5300e+00, 4.4670e+01,
3.4200e+00, 2.6640e+01, 1.6910e+01, 1.2900e+01, 6.7900e+00,
8.7600e+00, 2.1000e-01, 1.9680e+01, 1.4270e+01, 1.6190e+01,
5.0300e+00, 1.4100e+01, 1.0239e+02, 4.2600e+00, 2.7230e+01,
5.9180e+01, 1.3350e+01, 8.9640e+01, 8.0800e+01, 8.9500e+00,
1.2459e+02, 5.8000e+00, 2.1600e+01, 1.4597e+02, 7.6470e+01,
4.7600e+01, 2.8400e+01, 2.5970e+01, 8.5810e+01, 4.5370e+01,
1.4300e+00, 1.5900e+00, 8.2990e+01, 8.6700e+01, 1.5140e+01,
3.1380e+01, 7.1700e+00, 1.8240e+01, 5.8860e+01, 6.0600e+00,
2.2800e+01, 5.3700e+00, 7.7470e+01, 1.5200e+01, 6.8014e+02,
7.1000e-01, 1.4400e+00, 2.0000e-02, 4.6000e+00, 3.0850e+01,
8.2100e+00, 4.9310e+01, 3.6900e+00, 1.1250e+01, 3.0800e+00,
1.6030e+01, 3.4300e+00, 8.9000e-01, 1.2800e+00, 5.4200e+00,
5.4100e+00, 7.2700e+00, 8.6000e+00, 3.6600e+00, 3.2500e+00,
2.6280e+01, 2.2000e-01, 1.8200e+00, 8.2000e-01, 2.7990e+01,
2.4700e+00, 7.5400e+00, 2.7100e+01, 4.0000e-02, 6.7000e-01,
7.2000e-01, 8.4000e-01, 7.0000e-02, 1.2000e-01, 6.7300e+00,
5.7000e-01, 6.1400e+00, 1.1500e+00, 3.6290e+01, 3.1000e-01,
1.2200e+00, 2.3000e-01, 8.4437e+02, 1.9500e+00, 9.0000e-02,
7.4800e+00, 9.4500e+00, 2.7700e+00, 2.7280e+01, 7.3000e-01,
4.7000e-01, 1.7590e+02, 1.3540e+01, 4.1820e+01, 2.9760e+01,
3.0000e-02, 8.5130e+01, 1.0980e+01, 1.7302e+02, 5.2000e-01,
1.2400e+00, 1.1900e+00, 1.3000e-01, 1.0530e+01, 5.1000e-01,
1.5000e-01, 4.9000e-01, 7.9000e-01, 1.9450e+01, 4.4000e-01,
3.5000e-01, 3.5300e+00, 8.0000e-02, 4.3000e-01, 6.4000e-01,
3.7000e-01, 9.7000e-01, 1.9000e-01, 1.4000e-01, 3.5400e+00,
3.9200e+00, 3.3000e-01, 8.8400e+01, 2.0890e+01, 9.4750e+01,
8.3960e+01, 3.9731e+02, 1.6041e+02, 2.3970e+01, 7.0560e+01,
```

```
1.1330e+01, 1.3200e+00, 2.4000e-01, 6.5000e-01, 8.0000e-01,  
3.9000e-01, 1.3300e+01, 2.9000e-01, 3.3000e+00, 1.6090e+01,  
1.6600e+00, 7.0000e-01, 6.3000e-01, 9.3000e-01, 1.2000e+00,  
1.0300e+00, 2.8600e+00, 1.5500e+00, 2.3900e+00, 1.3000e+00,  
6.0000e-01, 1.7000e-01, 4.0000e-01, 3.6800e+00, 8.5050e+01,  
8.4700e+00, 2.5312e+02, 4.5000e-01, 6.4100e+00, 5.9000e-01,  
2.5000e-01, 6.5000e+00, 6.3100e+00, 1.6500e+00, 5.6100e+00,  
6.9000e-01, 9.9000e-01, 9.4300e+00, nan, 8.8340e+01,  
3.5700e+00, 1.8000e+00, 1.9100e+00, 2.4000e+00])
```



```
In [23]: trial['Money_Value'].unique()
```

```

Out[23]: array([3.3800e+00, 9.4000e-01, 0.0000e+00, 1.1750e+01, 2.9500e+00,
4.4950e+01, 7.7900e+00, 7.3400e+00, 1.9300e+00, 4.4200e+00,
9.6000e-01, 1.0430e+01, 7.0000e-03, 9.0000e+00, 4.1280e+01,
1.4030e+01, 6.3180e+01, 3.4240e+01, 1.0000e-02, 2.0519e+02,
1.0000e-01, 1.1160e+01, 1.2500e+00, 1.4600e+00, 6.7800e+00,
1.1600e+00, 1.5241e+02, 1.0800e+00, 2.8400e+00, 9.0000e-01,
9.6700e+00, 3.2680e+01, 9.3503e+02, 2.9630e+01, 1.1000e-01,
2.6200e+00, 6.0000e-02, 2.4300e+00, 1.2613e+02, 2.0790e+01,
1.5692e+02, 1.2290e+01, 2.2900e+00, 7.7800e+00, 2.5100e+00,
8.3100e+00, 1.6000e-01, 4.7900e+01, 8.9100e+00, 4.9500e+00,
1.7500e+00, 6.8000e-01, 1.5820e+01, 5.8000e-01, 2.1531e+02,
5.0000e-02, 5.3340e+01, 1.0690e+01, 5.6900e+00, 1.3500e+00,
1.1690e+01, 1.4600e+01, 2.0780e+01, 1.0222e+02, 7.4000e-01,
8.7337e+02, 1.4000e+00, 4.2000e-01, 3.6520e+01, 6.0200e+00,
1.7160e+01, 1.2910e+01, 1.0790e+01, 3.4600e+00, 2.3300e+00,
5.5800e+00, 7.6000e-01, 2.5270e+01, 2.0000e-01, 7.5600e+00,
6.7030e+01, 1.9400e+01, 3.4830e+01, 9.8750e+01, 3.6000e-01,
3.5210e+01, 8.2000e+00, 2.0330e+01, 3.5130e+01, 6.3700e+01,
2.8000e-01, 2.7130e+01, 1.0270e+01, 1.3050e+01, 1.8790e+01,
1.6820e+01, 1.0030e+01, 1.2670e+01, 1.3790e+01, 1.0160e+01,
9.7300e+00, 2.6950e+01, 2.9070e+01, 9.7600e+00, 1.3310e+01,
7.7000e+00, 1.0400e+01, 2.0620e+01, 4.4020e+01, 4.3530e+01,
1.2030e+01, 1.1880e+01, 6.9600e+00, 2.1070e+01, 5.7160e+01,
1.0650e+01, 5.3600e+00, 2.4500e+00, 3.1610e+01, 3.4320e+01,
4.2400e+00, 8.3800e+00, 2.2100e+01, 5.4070e+01, 2.7680e+01,
2.1450e+01, 1.1090e+01, 1.0140e+01, 5.2130e+01, 1.7020e+01,
1.8000e+01, 6.9760e+01, 1.3880e+01, 4.0300e+00, 2.4400e+00,
7.8980e+01, 6.8800e+00, 5.4900e+00, 3.7790e+01, 6.0880e+01,
1.8320e+02, 1.0020e+01, 3.2000e-01, 1.8000e-01, 4.8560e+01,
1.8450e+01, 8.4400e+00, 3.9650e+01, 9.5300e+00, 4.4670e+01,
3.4200e+00, 2.6640e+01, 1.6910e+01, 1.2900e+01, 6.7900e+00,
8.7600e+00, 2.1000e-01, 1.9680e+01, 1.4270e+01, 1.6190e+01,
5.0300e+00, 1.4100e+01, 1.0239e+02, 4.2600e+00, 2.7230e+01,
5.9180e+01, 1.3350e+01, 8.9640e+01, 8.0800e+01, 8.9500e+00,
1.2459e+02, 5.8000e+00, 2.1600e+01, 1.4597e+02, 7.6470e+01,
4.7600e+01, 2.8400e+01, 2.5970e+01, 8.5810e+01, 4.5370e+01,
1.4300e+00, 1.5900e+00, 8.2990e+01, 8.6700e+01, 1.5140e+01,
3.1380e+01, 7.1700e+00, 1.8240e+01, 5.8860e+01, 6.0600e+00,
2.2800e+01, 5.3700e+00, 7.7470e+01, 1.5200e+01, 6.8014e+02,
7.1000e-01, 1.4400e+00, 2.0000e-02, 4.6000e+00, 3.0850e+01,
8.2100e+00, 4.9310e+01, 3.6900e+00, 1.1250e+01, 3.0800e+00,
1.6030e+01, 3.4300e+00, 8.9000e-01, 1.2800e+00, 5.4200e+00,
5.4100e+00, 7.2700e+00, 8.6000e+00, 3.6600e+00, 3.2500e+00,
2.6280e+01, 2.2000e-01, 1.8200e+00, 8.2000e-01, 2.7990e+01,
2.4700e+00, 7.5400e+00, 2.7100e+01, 4.0000e-02, 6.7000e-01,
7.2000e-01, 8.4000e-01, 7.0000e-02, 1.2000e-01, 6.7300e+00,
5.7000e-01, 6.1400e+00, 1.1500e+00, 3.6290e+01, 3.1000e-01,
1.2200e+00, 2.3000e-01, 8.4437e+02, 1.9500e+00, 9.0000e-02,
7.4800e+00, 9.4500e+00, 2.7700e+00, 2.7280e+01, 7.3000e-01,
4.7000e-01, 1.7590e+02, 1.3540e+01, 4.1820e+01, 2.9760e+01,
3.0000e-02, 8.5130e+01, 1.0980e+01, 1.7302e+02, 5.2000e-01,
1.2400e+00, 1.1900e+00, 1.3000e-01, 1.0530e+01, 5.1000e-01,
1.5000e-01, 4.9000e-01, 7.9000e-01, 1.9450e+01, 4.4000e-01,
3.5000e-01, 3.5300e+00, 8.0000e-02, 4.3000e-01, 6.4000e-01,
3.7000e-01, 9.7000e-01, 1.9000e-01, 1.4000e-01, 3.5400e+00,
3.9200e+00, 3.3000e-01, 8.8400e+01, 2.0890e+01, 9.4750e+01,
8.3960e+01, 3.9731e+02, 1.6041e+02, 2.3970e+01, 7.0560e+01,

```

```
1.1330e+01, 1.3200e+00, 2.4000e-01, 6.5000e-01, 8.0000e-01,
3.9000e-01, 1.3300e+01, 2.9000e-01, 3.3000e+00, 1.6090e+01,
1.6600e+00, 7.0000e-01, 6.3000e-01, 9.3000e-01, 1.2000e+00,
1.0300e+00, 2.8600e+00, 1.5500e+00, 2.3900e+00, 1.3000e+00,
6.0000e-01, 1.7000e-01, 4.0000e-01, 3.6800e+00, 8.5050e+01,
8.4700e+00, 2.5312e+02, 4.5000e-01, 6.4100e+00, 5.9000e-01,
2.5000e-01, 6.5000e+00, 6.3100e+00, 1.6500e+00, 5.6100e+00,
6.9000e-01, 9.9000e-01, 9.4300e+00, nan, 8.8340e+01,
3.5700e+00, 1.8000e+00, 1.9100e+00, 2.4000e+00])
```

Both audit & trail have one missing value. Since dataset is small its better to impute the missing value rather than removing missing values. So let's impute the missing Money\_Value with it's mean grouped by 'numbers' column

```
In [24]: audit[audit['Money_Value'].isnull()]
```

Out[24]:

	Sector_score	LOCATION_ID	PARA_A	Score_A	Risk_A	PARA_B	Score_B	Risk_B	TOTAL
<b>642</b>	55.57	4	0.23	0.2	0.046	0.0	0.2	0.0	0.2

1 rows × 27 columns

```
In [25]: audit['Sector_score'].value_counts()
```

```
Out[25]: 55.57    200
          3.89    114
          1.85     95
          2.72     82
          3.41     76
          2.37     74
          1.99     47
          21.61    41
          59.85    37
          2.34      5
          15.56     3
          2.36      1
          17.68     1
Name: Sector_score, dtype: int64
```

```
In [26]: audit[['Sector_score', 'Money_Value']].groupby('Sector_score').mean()
```

Out[26]:

Money_Value	
Sector_score	
1.85	2.401579
1.99	26.892340
2.34	16.792000
2.36	88.400000
2.37	5.569189
2.72	30.059512
3.41	18.294737
3.89	33.157842
15.56	193.896667
17.68	70.560000
21.61	0.449024
55.57	2.574221
59.85	1.672703

```
In [27]: audit[['Sector_score', 'Money_Value']].groupby('Sector_score').median()
```

Out[27]:

Money_Value	
Sector_score	
1.85	0.020
1.99	0.050
2.34	0.000
2.36	88.400
2.37	0.575
2.72	8.325
3.41	10.215
3.89	2.565
15.56	160.410
17.68	70.560
21.61	0.000
55.57	0.000
59.85	0.290

Sector\_Score - Is the score of each firm mentioned above in Data Description

Missing value is in 'Money\_Value' that falls under Sector\_Score = 55.57.

Impute the missing value with median of Money\_Value under Sector\_score = 55.57

```
In [28]: audit['Money_Value'].fillna(0,inplace=True)
         trial['Money_Value'].fillna(0,inplace=True)
```

## 2. Merging Datasets

Both dataframes (audit, trial) have similar columns that hold similar value

First sort by common columns and concatenate later

So selecting only unique column from trial data frame and concatenating to audit for analysis

Removing the column ('LOCATION\_ID') on which analysis is not done

```
In [29]: # Sorting the data frames
         audit = audit.sort_values(by=['LOCATION_ID', 'TOTAL'])
         trial = trial.sort_values(by=['LOCATION_ID', 'TOTAL'])

         #Concatinate data frames
         data = pd.concat([audit,trial[['Loss', 'History_score', 'LOSS_SCORE', 'MONEY_Marks']]],axis=1)

         #Removing LOCATION_ID
         data.drop('LOCATION_ID',axis=1,inplace=True)
```

## 3. Handling Outliers

In data distribution section we can across an outlier with large Audit\_Risk, TOTAL value. We can handle this enither by imputing with a mean value or removing the outlier value.

Let's remove the outlier observation from data

```
In [30]: data['TOTAL'].max()
         # 1268.91
```

```
Out[30]: 1268.91
```

```
In [31]: data[data['TOTAL'] == 1268.91]
```

```
Out[31]:
```


	Sector_score	PARA_A	Score_A	Risk_A	PARA_B	Score_B	Risk_B	TOTAL	numbers	Score_B.1
241	2.72	4.28	0.6	2.568	1264.63	0.6	758.778	1268.91	5.5	0.6

1 rows × 30 columns

```
In [32]: # index = 241
data.drop(241,inplace=True)
```

```
In [33]: #looking for outliers in the risk colum after removing outlier
sns.stripplot(x='Risk',y='Audit_Risk',data=data)
```

```
Out[33]: <matplotlib.axes._subplots.AxesSubplot at 0x1e46aba3b00>
```



## Feature Selection


```
In [34]: data.columns
```

```
Out[34]: Index(['Sector_score', 'PARA_A', 'Score_A', 'Risk_A', 'PARA_B', 'Score_B',
               'Risk_B', 'TOTAL', 'numbers', 'Score_B.1', 'Risk_C', 'Money_Value',
               'Score_MV', 'Risk_D', 'District_Loss', 'PROB', 'Risk_E', 'History',
               'Prob', 'Risk_F', 'Score', 'Inherent_Risk', 'CONTROL_RISK',
               'Detection_Risk', 'Audit_Risk', 'Risk', 'Loss', 'History_score',
               'LOSS_SCORE', 'MONEY_Marks'],
              dtype='object')
```

```
In [35]: fig, ax = plt.subplots(figsize=(20, 20))
corr = data.corr()
colormap = sns.diverging_palette(220, 10, as_cmap=True)
sns.heatmap(corr, cmap=colormap, annot=True, fmt=".2f")
plt.xticks(range(len(corr.columns)), corr.columns)
plt.yticks(range(len(corr.columns)), corr.columns)
```

```
Out[35]: ([<matplotlib.axis.YTick at 0x1e46abfab70>,  
<matplotlib.axis.YTick at 0x1e46abfa438>,  
<matplotlib.axis.YTick at 0x1e46abf9390>,  
<matplotlib.axis.YTick at 0x1e46b4c3978>,  
<matplotlib.axis.YTick at 0x1e46b4c3e48>,  
<matplotlib.axis.YTick at 0x1e46b4cd358>,  
<matplotlib.axis.YTick at 0x1e46b4cd828>,  
<matplotlib.axis.YTick at 0x1e46b4cdcfc8>,  
<matplotlib.axis.YTick at 0x1e46b4d7208>,  
<matplotlib.axis.YTick at 0x1e46b4d7710>,  
<matplotlib.axis.YTick at 0x1e46b4d7c18>,  
<matplotlib.axis.YTick at 0x1e46b4d7cf8>,  
<matplotlib.axis.YTick at 0x1e46b4cd780>,  
<matplotlib.axis.YTick at 0x1e46b4de160>,  
<matplotlib.axis.YTick at 0x1e46b4de5f8>,  
<matplotlib.axis.YTick at 0x1e46b4deb00>,  
<matplotlib.axis.YTick at 0x1e46b4e70b8>,  
<matplotlib.axis.YTick at 0x1e46b4e7550>,  
<matplotlib.axis.YTick at 0x1e46b4e7a58>,  
<matplotlib.axis.YTick at 0x1e46b4e7f60>,  
<matplotlib.axis.YTick at 0x1e46b4e75f8>,  
<matplotlib.axis.YTick at 0x1e46b4deac8>,  
<matplotlib.axis.YTick at 0x1e46b4f04a8>,  
<matplotlib.axis.YTick at 0x1e46b4f08d0>,  
<matplotlib.axis.YTick at 0x1e46b4f0dd8>,  
<matplotlib.axis.YTick at 0x1e46b4f7320>,  
<matplotlib.axis.YTick at 0x1e46b4f7828>,  
<matplotlib.axis.YTick at 0x1e46b4f7d30>,  
<matplotlib.axis.YTick at 0x1e46b4fe278>,  
<matplotlib.axis.YTick at 0x1e46b4f77f0>],  
<a list of 30 Text yticklabel objects>)
```





Highly correlated variables deflects model's accuracy. So it's better to remove variables that are highly correlated. Choosing correlation = 0.7 as threshold and removing variables that have correlation greater than 0.7

Below are the columns that that are considered for analysis which are less correlated

Removing Risk column since it's formed due to Audit\_Risk score. Being Risk(1/0) is due to Audit\_Risk.

```
In [36]: data = data[['Sector_score', 'PARA_A', 'Score_A', 'PARA_B', 'Score_B', 'numbers',
                    'Money_Value', 'Score_MV', 'District_Loss', 'LOSS_SCORE', 'History_score', 'Audit_Risk']]
```

```
In [37]: data.columns
```

```
Out[37]: Index(['Sector_score', 'PARA_A', 'Score_A', 'PARA_B', 'Score_B', 'numbers',  
              'Money_Value', 'Score_MV', 'District_Loss', 'LOSS_SCORE',  
              'History_score', 'Audit_Risk'],  
            dtype='object')
```

## Data Scaling

Since the variables we considered for analysis are not in same range we need to scale them before analysis.

```
In [38]: X = data.drop('Audit_Risk',axis=1)  
        y = data['Audit_Risk']  
  
        from sklearn.model_selection import train_test_split  
        X_train_org, X_test_org, y_train, y_test = train_test_split(X, y, test_size=0.  
        2, random_state=0)  
  
        from sklearn.preprocessing import MinMaxScaler  
        scaler = MinMaxScaler()  
        X_train = scaler.fit_transform(X_train_org)  
        X_test = scaler.transform(X_test_org)
```

```
C:\Users\tanmay\Anaconda3\lib\site-packages\sklearn\preprocessing\data.py:32  
3: DataConversionWarning: Data with input dtype int64, float64 were all converted to float64 by MinMaxScaler.  
    return self.partial_fit(X, y)
```

## Supervised Learning - Regression

1. Linear Regression
2. KNN Regressor
3. Ridge Regression
4. Lasso Regression
5. Polynomial Regression
6. SVM (simple & kernel)

### 1.Linear Regression

```
In [111]: #Linear Regression
from sklearn import metrics
from sklearn.linear_model import LinearRegression
from sklearn.metrics import accuracy_score
lr = LinearRegression()
lr.fit(X_train,y_train)
print("Train:%.4f"%lr.score(X_train,y_train))
print("Test:%.4f"%lr.score(X_test,y_test))
pred = lr.predict(X_test)

print("Mean Squared Error Test:", metrics.mean_squared_error(y_test,pred))
```

Train:0.7544

Test:0.7233

Mean Squared Error Test: 0.06791434815183471

```
In [40]: # Using GridSearch
from sklearn.model_selection import GridSearchCV
model = LinearRegression()
parameters = {'fit_intercept':[True,False], 'normalize':[True,False], 'copy_X':
:[True, False]}
grid = GridSearchCV(model,param_grid=parameters,cv=5)
grid.fit(X_train,y_train)
print("Train:%.4f"%grid.score(X_train,y_train))
print("Test:%.4f"%grid.score(X_test,y_test))
print("Best Parameters:{}",format(grid.best_params_))
```

Train:0.8002

Test:0.6659

Best Parameters:{} {'copy\_X': True, 'fit\_intercept': False, 'normalize': True}

```
In [41]: #cross validation
from sklearn.model_selection import cross_val_score
lr = LinearRegression()
train_score = cross_val_score(lr,X_train,y_train,cv=10)
test_score = cross_val_score(lr,X_test,y_test,cv=10)

print("Avg Train Score:%.4f"%train_score.mean())
print("Avg Test Score:%.4f"%test_score.mean())
```

Avg Train Score:0.7300

Avg Test Score:-0.0222

## 2. KNN Regression

```
In [113]: from sklearn.neighbors import KNeighborsRegressor
from sklearn.preprocessing import StandardScaler

X_train_org, X_test_org, y_train, y_test = train_test_split(X,y,test_size=0.2,
random_state = 0)

scaler = StandardScaler()
X_train = scaler.fit_transform(X_train_org)
X_test = scaler.transform(X_test_org)

train_score_array = []
test_score_array = []


for k in range(1,20):
    knn_reg = KNeighborsRegressor(k)
    knn_reg.fit(X_train, y_train)
    train_score_array.append(knn_reg.score(X_train, y_train))
    test_score_array.append(knn_reg.score(X_test, y_test))

print("Train score: {}".format(train_score_array))
print("Test score: {}".format(test_score_array))
x_axis = range(1,20)
plt.plot(x_axis, train_score_array, c = 'g', label = 'Train Score')
plt.plot(x_axis, test_score_array, c = 'b', label = 'Test Score')
plt.legend()
plt.xlabel('k')
plt.ylabel('MSE')
```

```
C:\Users\Tanmay\Anaconda3\lib\site-packages\sklearn\preprocessing\data.py:62
5: DataConversionWarning: Data with input dtype int64, float64 were all converted to float64 by StandardScaler.
    return self.partial_fit(X, y)
C:\Users\Tanmay\Anaconda3\lib\site-packages\sklearn\base.py:462: DataConversionWarning: Data with input dtype int64, float64 were all converted to float64 by StandardScaler.
    return self.fit(X, **fit_params).transform(X)
C:\Users\Tanmay\Anaconda3\lib\site-packages\ipykernel_launcher.py:8: DataConversionWarning: Data with input dtype int64, float64 were all converted to float64 by StandardScaler.
```

```
Train score: {} [1.0, 0.979508873979575, 0.95901774795915, 0.949625981866455
1, 0.9434444921836269, 0.9366292954553522, 0.9265386162397008, 0.923051552808
7164, 0.9195533571049981, 0.9169426358638773, 0.9112615699611347, 0.904943943
1830284, 0.902960170463431, 0.8982761958271758, 0.8939318602735332, 0.8906873
134040607, 0.8862234363756331, 0.8852328291613231, 0.8835241257786367]
Test score: {} [0.8948439620081411, 0.8619827001356852, 0.8568709482888588,
0.855410447761194, 0.8559362279511534, 0.842996193276044, 0.8428024534101292,
0.8373367537313433, 0.8322047171549659, 0.8222862957937584, 0.809676261816387
6, 0.8026498661993064, 0.8013546442076867, 0.8037713219616205, 0.802072968490
879, 0.8052970234056988, 0.8036056349269695, 0.8070516525118515, 0.8025775491
71794]
```

Out[113]: Text(0, 0.5, 'MSE')



```
In [114]: knn = KNeighborsRegressor(3)
knn.fit(X_train,y_train)
print(knn.score(X_train,y_train))
print(knn.score(X_test,y_test))
```

```
0.95901774795915
0.8568709482888588
```

```
In [115]: # Using Cross Validation
from sklearn.model_selection import cross_val_score
knn_reg = KNeighborsRegressor(n_neighbors=3)
train_score = cross_val_score(knn_reg,X_train,y_train,cv=10)
test_score = cross_val_score(knn_reg,X_test,y_test,cv=10)

print("Avg Train Score:%.4f"%train_score.mean())
print("Avg Test Score:%.4f"%test_score.mean())
pred = knn.predict(X_test)
print("Mean Squared Error Test:", metrics.mean_squared_error(y_test,pred))
```

Avg Train Score:0.9100

Avg Test Score:0.7501

Mean Squared Error Test: 0.03512544802867384

```


In [116]: k_range = range(1, 20)
# list of scores from k_range
k_train_scores = []
k_test_scores = []

# 1. we will loop through reasonable values of k
for k in k_range:
    # 2. run KNeighborsRegressor with k neighbours
    knn = KNeighborsRegressor(n_neighbors=k)
    # 3. obtain cross_val_score for KNeighborsRegressor with k neighbours
    scores1 = cross_val_score(knn, X_train, y_train, cv=10)
    scores2 = cross_val_score(knn, X_test, y_test, cv=10)
    # 4. append mean of scores for k neighbors to k_scores list
    k_train_scores.append(scores1.mean())
    k_test_scores.append(scores2.mean())

# plot the value of K for KNN (x-axis) versus the cross-validated accuracy (y-axis)
plt.plot(k_range, k_train_scores, 'b', label='Train Score')
plt.plot(k_range, k_test_scores, 'r', label='Test Score')
plt.legend()
plt.xlabel('Value of K for KNN')
plt.ylabel('Cross-Validated Accuracy')

```

Out[116]: Text(0, 0.5, 'Cross-Validated Accuracy')



```
In [118]: #Most efficient parameter using GridSearch()
from sklearn.model_selection import GridSearchCV
k_range = list(range(1, 31))
param_grid = dict(n_neighbors=k_range)
grid = GridSearchCV(knn, param_grid, cv=10)
grid.fit(X_train,y_train)
grid.best_score_
grid.best_estimator_

print("Best parameters: {}".format(grid.best_params_))
print("Best cross-validation score: {:.2f}".format(grid.best_score_))

results = pd.DataFrame(grid.cv_results_)
# show the first 5 rows
display(results.head())
```



```
Best parameters: {'n_neighbors': 1}  
Best cross-validation score: 0.91
```

```
C:\Users\tanmay\Anaconda3\lib\site-packages\sklearn\utils\deprecation.py:125:
FutureWarning: You are accessing a training score ('split0_train_score'), whi
ch will not be available by default any more in 0.21. If you need training sc
ores, please set return_train_score=True
  warnings.warn(*warn_args, **warn_kwargs)
C:\Users\tanmay\Anaconda3\lib\site-packages\sklearn\utils\deprecation.py:125:
FutureWarning: You are accessing a training score ('split1_train_score'), whi
ch will not be available by default any more in 0.21. If you need training sc
ores, please set return_train_score=True
  warnings.warn(*warn_args, **warn_kwargs)
C:\Users\tanmay\Anaconda3\lib\site-packages\sklearn\utils\deprecation.py:125:
FutureWarning: You are accessing a training score ('split2_train_score'), whi
ch will not be available by default any more in 0.21. If you need training sc
ores, please set return_train_score=True
  warnings.warn(*warn_args, **warn_kwargs)
C:\Users\tanmay\Anaconda3\lib\site-packages\sklearn\utils\deprecation.py:125:
FutureWarning: You are accessing a training score ('split3_train_score'), whi
ch will not be available by default any more in 0.21. If you need training sc
ores, please set return_train_score=True
  warnings.warn(*warn_args, **warn_kwargs)
C:\Users\tanmay\Anaconda3\lib\site-packages\sklearn\utils\deprecation.py:125:
FutureWarning: You are accessing a training score ('split4_train_score'), whi
ch will not be available by default any more in 0.21. If you need training sc
ores, please set return_train_score=True
  warnings.warn(*warn_args, **warn_kwargs)
C:\Users\tanmay\Anaconda3\lib\site-packages\sklearn\utils\deprecation.py:125:
FutureWarning: You are accessing a training score ('split5_train_score'), whi
ch will not be available by default any more in 0.21. If you need training sc
ores, please set return_train_score=True
  warnings.warn(*warn_args, **warn_kwargs)
C:\Users\tanmay\Anaconda3\lib\site-packages\sklearn\utils\deprecation.py:125:
FutureWarning: You are accessing a training score ('split6_train_score'), whi
ch will not be available by default any more in 0.21. If you need training sc
ores, please set return_train_score=True
  warnings.warn(*warn_args, **warn_kwargs)
C:\Users\tanmay\Anaconda3\lib\site-packages\sklearn\utils\deprecation.py:125:
FutureWarning: You are accessing a training score ('split7_train_score'), whi
ch will not be available by default any more in 0.21. If you need training sc
ores, please set return_train_score=True
  warnings.warn(*warn_args, **warn_kwargs)
C:\Users\tanmay\Anaconda3\lib\site-packages\sklearn\utils\deprecation.py:125:
FutureWarning: You are accessing a training score ('split8_train_score'), whi
ch will not be available by default any more in 0.21. If you need training sc
ores, please set return_train_score=True
  warnings.warn(*warn_args, **warn_kwargs)
C:\Users\tanmay\Anaconda3\lib\site-packages\sklearn\utils\deprecation.py:125:
FutureWarning: You are accessing a training score ('split9_train_score'), whi
ch will not be available by default any more in 0.21. If you need training sc
ores, please set return_train_score=True
  warnings.warn(*warn_args, **warn_kwargs)
C:\Users\tanmay\Anaconda3\lib\site-packages\sklearn\utils\deprecation.py:125:
FutureWarning: You are accessing a training score ('mean_train_score'), whic
h will not be available by default any more in 0.21. If you need training sco
res, please set return_train_score=True
  warnings.warn(*warn_args, **warn_kwargs)
C:\Users\tanmay\Anaconda3\lib\site-packages\sklearn\utils\deprecation.py:125:
FutureWarning: You are accessing a training score ('std_train_score'), which
```

will not be available by default any more in 0.21. If you need training score  
 s, please set return\_train\_score=True  
 warnings.warn(\*warn\_args, \*\*warn\_kwargs)

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_n_neighbors	parar
<b>0</b>	0.000997	0.000446	0.001297	0.000457	1	{'n_neighbor
<b>1</b>	0.000898	0.000299	0.001097	0.000299	2	{'n_neighbor
<b>2</b>	0.000997	0.000446	0.001097	0.000299	3	{'n_neighbor
<b>3</b>	0.001094	0.000538	0.001299	0.000455	4	{'n_neighbor
<b>4</b>	0.000693	0.000454	0.001388	0.000481	5	{'n_neighbor

5 rows × 31 columns

### 3.Ridge Regression


```
In [120]: from sklearn.linear_model import Ridge

x_range = [0.01, 0.1, 1, 10, 100, 1000]
train_score_list = []
test_score_list = []

for alpha in x_range:
    ridge = Ridge(alpha)
    ridge.fit(X_train, y_train)
    train_score_list.append(ridge.score(X_train, y_train))
    test_score_list.append(ridge.score(X_test, y_test))

plt.plot(x_range, train_score_list, c = 'g', label = 'Train Score')
plt.plot(x_range, test_score_list, c = 'b', label = 'Test Score')
plt.xscale('log')
plt.legend(loc = 3)
plt.xlabel(r'$\alpha$')
```

Out[120]: Text(0.5, 0, '\$\alpha\$')



```
In [121]: ridge = Ridge(alpha = 0.01)
ridge.fit(X_train, y_train)
print('Train score: {:.4f}'.format(ridge.score(X_train, y_train)))
print('Test score: {:.4f}'.format(ridge.score(X_test, y_test)))
pred = ridge.predict(X_test)
print("Mean Squared Error Test:", metrics.mean_squared_error(y_test, pred))
```

Train score: 0.7544

Test score: 0.7233

Mean Squared Error Test: 0.0679143177607114

```

In [49]: x_range1 = np.linspace(0.001, 1, 100).reshape(-1,1)
x_range2 = np.linspace(1, 10000, 100).reshape(-1,1)

x_range = np.append(x_range1, x_range2)
coeff = []

for alpha in x_range:
    ridge = Ridge(alpha)
    ridge.fit(X_train,y_train)
    coeff.append(ridge.coef_)

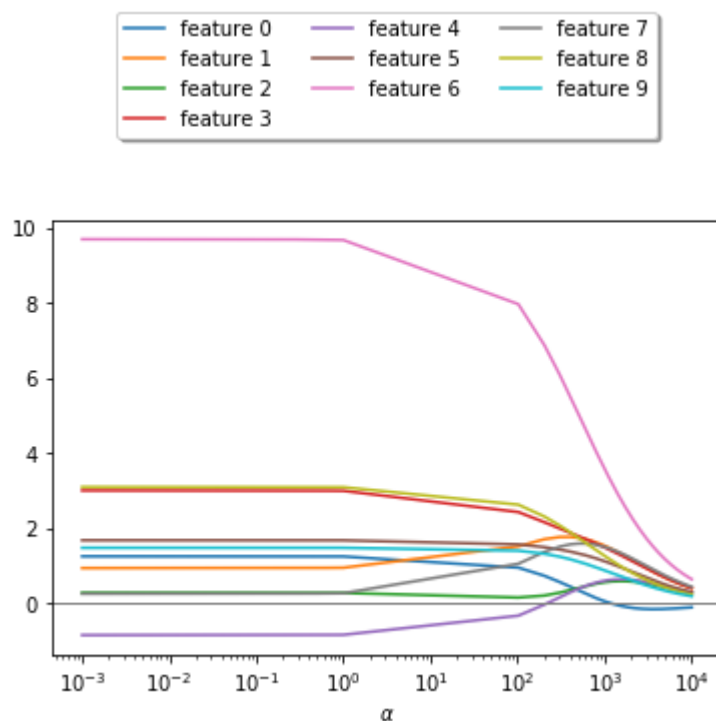
coeff = np.array(coeff)

for i in range(0,10):
    plt.plot(x_range, coeff[:,i], label = 'feature {:d}'.format(i))

plt.axhline(y=0, xmin=0.001, xmax=9999, linewidth=1, c = 'gray')
plt.xlabel(r'$\alpha$')
plt.xscale('log')
plt.legend(loc='upper center', bbox_to_anchor=(0.5, 1.5),
          ncol=3, fancybox=True, shadow=True)

```

Out[49]: <matplotlib.legend.Legend at 0x1e46c262e10>



```
In [125]: #cross validation
ridge = Ridge(alpha = 1000)
train_score = cross_val_score(ridge,X_train,y_train,cv=5)
test_score = cross_val_score(ridge,X_test,y_test,cv=5)
print("Avg Train Score:%.4f"%train_score.mean())
print("Avg Test Score:%.4f"%test_score.mean())
```

Avg Train Score:0.5754

Avg Test Score:0.3051

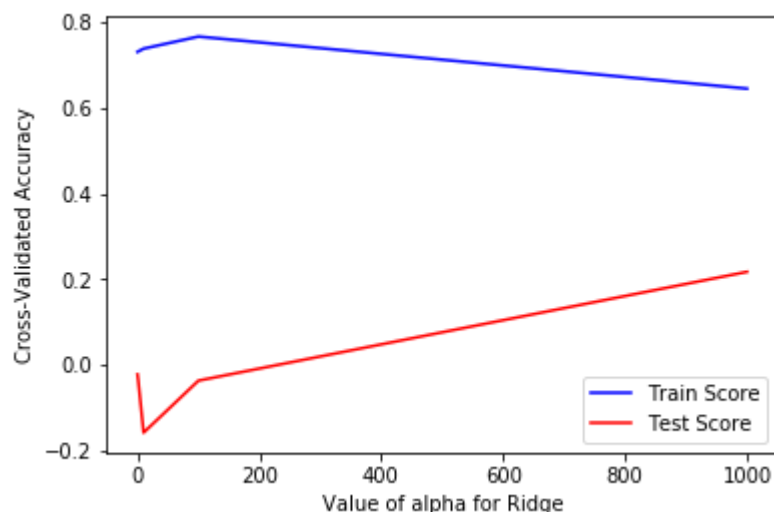
```
In [51]: alpha = [0.001,0.01, 0.1,1,10,100,1000]

ridge_train_scores = []
ridge_test_scores = []

# 1. we will loop through reasonable values of k
for i in alpha:
    # 2. run KNeighborsClassifier with k neighbours
    ridge = Ridge(i)
    # 3. obtain cross_val_score for KNeighborsClassifier with k neighbours
    scores1 = cross_val_score(ridge, X_train, y_train, cv=10)
    scores2 = cross_val_score(ridge, X_test, y_test, cv=10)
    # 4. append mean of scores for k neighbors to k_scores list
    ridge_train_scores.append(scores1.mean())
    ridge_test_scores.append(scores2.mean())

# plot the value of K for KNN (x-axis) versus the cross-validated accuracy (y-axis)
plt.plot(alpha, ridge_train_scores,'b',label='Train Score')
plt.plot(alpha, ridge_test_scores,'r',label='Test Score')
plt.legend()
plt.xlabel('Value of alpha for Ridge')
plt.ylabel('Cross-Validated Accuracy')
```

Out[51]: Text(0, 0.5, 'Cross-Validated Accuracy')



```
In [52]: # GridSearch()
from sklearn.model_selection import GridSearchCV
alphas = np.array([1000,100,10,1,0.1,0.01,0.001])
model = Ridge()
grid = GridSearchCV(estimator=model, param_grid=dict(alpha=alphas))
grid.fit(X_train,y_train)

print("Best parameters: {}".format(grid.best_params_))
print("Best cross-validation score: {:.2f}".format(grid.best_score_))

results = pd.DataFrame(grid.cv_results_)
# show the first 5 rows
display(results.head())
```

Best parameters: {'alpha': 100.0}

Best cross-validation score: 0.75

C:\Users\tanmay\Anaconda3\lib\site-packages\sklearn\model\_selection\\_split.py:2053: FutureWarning: You should specify a value for 'cv' instead of relying on the default value. The default value will change from 3 to 5 in version 0.22.

warnings.warn(CV\_WARNING, FutureWarning)

C:\Users\tanmay\Anaconda3\lib\site-packages\sklearn\model\_selection\\_search.py:841: DeprecationWarning: The default of the 'iid' parameter will change from True to False in version 0.22 and will be removed in 0.24. This will change numeric results when test-set sizes are unequal.

DeprecationWarning)

C:\Users\tanmay\Anaconda3\lib\site-packages\sklearn\utils\deprecation.py:125: FutureWarning: You are accessing a training score ('split0\_train\_score'), which will not be available by default any more in 0.21. If you need training scores, please set return\_train\_score=True

warnings.warn(\*warn\_args, \*\*warn\_kwargs)

C:\Users\tanmay\Anaconda3\lib\site-packages\sklearn\utils\deprecation.py:125: FutureWarning: You are accessing a training score ('split1\_train\_score'), which will not be available by default any more in 0.21. If you need training scores, please set return\_train\_score=True

warnings.warn(\*warn\_args, \*\*warn\_kwargs)

C:\Users\tanmay\Anaconda3\lib\site-packages\sklearn\utils\deprecation.py:125: FutureWarning: You are accessing a training score ('split2\_train\_score'), which will not be available by default any more in 0.21. If you need training scores, please set return\_train\_score=True

warnings.warn(\*warn\_args, \*\*warn\_kwargs)

C:\Users\tanmay\Anaconda3\lib\site-packages\sklearn\utils\deprecation.py:125: FutureWarning: You are accessing a training score ('mean\_train\_score'), which will not be available by default any more in 0.21. If you need training scores, please set return\_train\_score=True

warnings.warn(\*warn\_args, \*\*warn\_kwargs)

C:\Users\tanmay\Anaconda3\lib\site-packages\sklearn\utils\deprecation.py:125: FutureWarning: You are accessing a training score ('std\_train\_score'), which will not be available by default any more in 0.21. If you need training scores, please set return\_train\_score=True

warnings.warn(\*warn\_args, \*\*warn\_kwargs)

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_alpha	params	split0_te
0	0.000665	4.703589e-04	0.000332	0.00047	1000	{'alpha': 1000.0}	
1	0.000665	4.701341e-04	0.000332	0.00047	100	{'alpha': 100.0}	
2	0.000665	4.703027e-04	0.000332	0.00047	10	{'alpha': 10.0}	
3	0.000997	6.257699e-07	0.000000	0.000000	1	{'alpha': 1.0}	
4	0.000665	4.701903e-04	0.000332	0.00047	0.1	{'alpha': 0.1}	




## 4.Lasso

```
In [126]: from sklearn.linear_model import Lasso
x_range = [0.01, 0.1, 1, 10, 100]
train_score_list = []
test_score_list = []

for alpha in x_range:
    lasso = Lasso(alpha)
    lasso.fit(X_train,y_train)
    train_score_list.append(lasso.score(X_train,y_train))
    test_score_list.append(lasso.score(X_test, y_test))

plt.plot(x_range, train_score_list, c = 'g', label = 'Train Score')
plt.plot(x_range, test_score_list, c = 'b', label = 'Test Score')
plt.xscale('log')
plt.legend(loc = 3)
plt.xlabel(r'$\alpha$')
```

Out[126]: Text(0.5, 0, '\$\alpha\$')



```
In [128]: lasso = Lasso(alpha=0.01)
lasso.fit(X_train,y_train)
print('Train score: {:.4f}'.format(lasso.score(X_train,y_train)))
print('Test score: {:.4f}'.format(lasso.score(X_test, y_test)))
pred = lasso.predict(X_test)
print("Mean Squared Error Test:", metrics.mean_squared_error(y_test,pred))
```

Train score: 0.7507

Test score: 0.7214

Mean Squared Error Test: 0.06837231408835201

```

In [55]: x_range1 = np.linspace(0.01, 1, 100).reshape(-1,1)
x_range2 = np.linspace(1, 100, 100).reshape(-1,1)

x_range = np.append(x_range1, x_range2)
coeff = []

for alpha in x_range:
    lasso = Lasso(alpha)
    lasso.fit(X_train,y_train)
    coeff.append(lasso.coef_)


coeff = np.array(coeff)

for i in range(0,10):
    plt.plot(x_range, coeff[:,i], label = 'feature {:d}'.format(i))

plt.axhline(y=0, xmin=0.001, xmax=9999, linewidth=1, c = 'gray')
plt.xlabel(r'$\alpha$')
plt.xscale('log')
plt.legend(loc='upper center', bbox_to_anchor=(0.5, 1.5),
          ncol=3, fancybox=True, shadow=True)

```

Out[55]: <matplotlib.legend.Legend at 0x1e46d8dbcf8>



```
In [56]: #Grid_Search
param_grid = {'alpha':[0.001, 0.01, 0.1, 1, 10, 100]}
grid_search = GridSearchCV(Lasso(),param_grid,cv=5,return_train_score=True)
grid_search.fit(X_train,y_train)

print("Best parameters: {}".format(grid_search.best_params_))
print("Best cross-validation score: {:.2f}".format(grid_search.best_score_))

results = pd.DataFrame(grid_search.cv_results_)
# show the first 5 rows
display(results.head())
```

Best parameters: {'alpha': 1}  
 Best cross-validation score: 0.78

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_alpha	params	split0_te
0	0.001197	0.000399	0.000599	0.000489	0.001	{'alpha': 0.001}	
1	0.000598	0.000488	0.000200	0.000399	0.01	{'alpha': 0.01}	
2	0.000798	0.000399	0.000200	0.000399	0.1	{'alpha': 0.1}	
3	0.000399	0.000489	0.000399	0.000489	1	{'alpha': 1}	
4	0.000399	0.000489	0.000199	0.000399	10	{'alpha': 10}	-

5 rows × 21 columns

```
In [57]: #cross validation
lasso = Lasso(alpha =0.001)
train_score = cross_val_score(lasso,X_train,y_train,cv=5)
test_score = cross_val_score(lasso,X_test,y_test,cv=5)
print("Avg Train Score: %.4f"%train_score.mean())
print("Avg Test Score: %.4f"%test_score.mean())
```

Avg Train Score:0.7572  
 Avg Test Score:-0.1782

## 5. Polynomial

```
In [138]: from sklearn.preprocessing import PolynomialFeatures

train_score_list = []
test_score_list = []

for n in range(1,3):
    poly = PolynomialFeatures(n)
    X_train_poly = poly.fit_transform(X_train)
    X_test_poly = poly.transform(X_test)
    lr.fit(X_train_poly, y_train)
    train_score_list.append(lr.score(X_train_poly, y_train))
    test_score_list.append(lr.score(X_test_poly, y_test))

print(train_score_list)
print(test_score_list)
pred = lr.predict(X_test_poly)
print("Mean Squared Error Test:", metrics.mean_squared_error(y_test,pred))

[0.7543570677111144, 0.9052075426722959]
[0.723262853740192, 0.692590358908533]
Mean Squared Error Test: 0.0754417167065677
```

## 6. SVM(simple)

```
In [59]: from sklearn.svm import LinearSVR
train_score = []
test_score = []
C = [0.01,0.1,1,10,100]

for i in C:
    svr = LinearSVR(C=i)
    svr.fit(X_train,y_train)
    train_score.append(svr.score(X_train,y_train))
    test_score.append(svr.score(X_test,y_test))


plt.plot(C,train_score,'b')
plt.plot(C,test_score,'r')
```

C:\Users\tanmay\Anaconda3\lib\site-packages\sklearn\svm\base.py:922: ConvergenceWarning: Liblinear failed to converge, increase the number of iterations.  
"the number of iterations.", ConvergenceWarning)

C:\Users\tanmay\Anaconda3\lib\site-packages\sklearn\svm\base.py:922: ConvergenceWarning: Liblinear failed to converge, increase the number of iterations.  
"the number of iterations.", ConvergenceWarning)

C:\Users\tanmay\Anaconda3\lib\site-packages\sklearn\svm\base.py:922: ConvergenceWarning: Liblinear failed to converge, increase the number of iterations.  
"the number of iterations.", ConvergenceWarning)

Out[59]: [<matplotlib.lines.Line2D at 0x1e46d97d9e8>]



```
In [60]: svr = LinearSVR(C=10)
svr.fit(X_train,y_train)
print('Train score: {:.4f}'.format(svr.score(X_train,y_train)))
print('Test score: {:.4f}'.format(svr.score(X_test, y_test)))
```

Train score: 0.7283  
Test score: 0.6079

C:\Users\tanmay\Anaconda3\lib\site-packages\sklearn\svm\base.py:922: ConvergenceWarning: Liblinear failed to converge, increase the number of iterations.  
"the number of iterations.", ConvergenceWarning)

```
In [61]: # Grid_Search
from sklearn.model_selection import GridSearchCV
param_grid = {'C':[0.001, 0.01, 0.1, 1, 10]}
grid_search = GridSearchCV(LinearSVR(),param_grid,cv=5,return_train_score=True
)

grid_search.fit(X_train,y_train)

print("Best parameters: {}".format(grid_search.best_params_))
print("Best cross-validation score: {:.2f}".format(grid_search.best_score_))

results = pd.DataFrame(grid_search.cv_results_)
# show the first 5 rows
display(results.head())
```

```

C:\Users\Tanmay\Anaconda3\lib\site-packages\sklearn\svm\base.py:922: Converge
nceWarning: Liblinear failed to converge, increase the number of iterations.
    "the number of iterations.", ConvergenceWarning)
C:\Users\Tanmay\Anaconda3\lib\site-packages\sklearn\svm\base.py:922: Converge
nceWarning: Liblinear failed to converge, increase the number of iterations.
    "the number of iterations.", ConvergenceWarning)
C:\Users\Tanmay\Anaconda3\lib\site-packages\sklearn\svm\base.py:922: Converge
nceWarning: Liblinear failed to converge, increase the number of iterations.
    "the number of iterations.", ConvergenceWarning)
C:\Users\Tanmay\Anaconda3\lib\site-packages\sklearn\svm\base.py:922: Converge
nceWarning: Liblinear failed to converge, increase the number of iterations.
    "the number of iterations.", ConvergenceWarning)
C:\Users\Tanmay\Anaconda3\lib\site-packages\sklearn\svm\base.py:922: Converge
nceWarning: Liblinear failed to converge, increase the number of iterations.
    "the number of iterations.", ConvergenceWarning)
C:\Users\Tanmay\Anaconda3\lib\site-packages\sklearn\svm\base.py:922: Converge
nceWarning: Liblinear failed to converge, increase the number of iterations.
    "the number of iterations.", ConvergenceWarning)
C:\Users\Tanmay\Anaconda3\lib\site-packages\sklearn\svm\base.py:922: Converge
nceWarning: Liblinear failed to converge, increase the number of iterations.
    "the number of iterations.", ConvergenceWarning)
C:\Users\Tanmay\Anaconda3\lib\site-packages\sklearn\svm\base.py:922: Converge
nceWarning: Liblinear failed to converge, increase the number of iterations.
    "the number of iterations.", ConvergenceWarning)
C:\Users\Tanmay\Anaconda3\lib\site-packages\sklearn\svm\base.py:922: Converge
nceWarning: Liblinear failed to converge, increase the number of iterations.
    "the number of iterations.", ConvergenceWarning)

```

Best parameters: {'C': 10}

Best cross-validation score: 0.77

```

C:\Users\Tanmay\Anaconda3\lib\site-packages\sklearn\svm\base.py:922: Converge
nceWarning: Liblinear failed to converge, increase the number of iterations.
    "the number of iterations.", ConvergenceWarning)

```

	mean_fit_time	std_fit_time	mean_score_time	std_score_time	param_C	params	split0_test_s
<b>0</b>	0.000599	0.000489	0.000399	0.000488	0.001	{'C': 0.001}	-0.09
<b>1</b>	0.001005	0.000016	0.000000	0.000000	0.01	{'C': 0.01}	0.08
<b>2</b>	0.001590	0.000483	0.000199	0.000399	0.1	{'C': 0.1}	0.51
<b>3</b>	0.015167	0.003713	0.000399	0.000489	1	{'C': 1}	0.62
<b>4</b>	0.023535	0.001493	0.000199	0.000399	10	{'C': 10}	0.62

5 rows × 21 columns

```
In [62]: #cross validation
svr = LinearSVR(C=10)
train_score = cross_val_score(svr,X_train,y_train,cv=5)
test_score = cross_val_score(svr,X_test,y_test,cv=5)
print("Avg Train Score: %.4f"%train_score.mean())
print("Avg Test Score: %.4f"%test_score.mean())
```

```
C:\Users\Tanmay\Anaconda3\lib\site-packages\sklearn\svm\base.py:922: Converge
nceWarning: Liblinear failed to converge, increase the number of iterations.
    "the number of iterations.", ConvergenceWarning)
C:\Users\Tanmay\Anaconda3\lib\site-packages\sklearn\svm\base.py:922: Converge
nceWarning: Liblinear failed to converge, increase the number of iterations.
    "the number of iterations.", ConvergenceWarning)
C:\Users\Tanmay\Anaconda3\lib\site-packages\sklearn\svm\base.py:922: Converge
nceWarning: Liblinear failed to converge, increase the number of iterations.
    "the number of iterations.", ConvergenceWarning)
C:\Users\Tanmay\Anaconda3\lib\site-packages\sklearn\svm\base.py:922: Converge
nceWarning: Liblinear failed to converge, increase the number of iterations.
    "the number of iterations.", ConvergenceWarning)
C:\Users\Tanmay\Anaconda3\lib\site-packages\sklearn\svm\base.py:922: Converge
nceWarning: Liblinear failed to converge, increase the number of iterations.
    "the number of iterations.", ConvergenceWarning)
```

Avg Train Score:0.7734

Avg Test Score:0.5096

```
C:\Users\Tanmay\Anaconda3\lib\site-packages\sklearn\svm\base.py:922: Converge
nceWarning: Liblinear failed to converge, increase the number of iterations.
    "the number of iterations.", ConvergenceWarning)
C:\Users\Tanmay\Anaconda3\lib\site-packages\sklearn\svm\base.py:922: Converge
nceWarning: Liblinear failed to converge, increase the number of iterations.
    "the number of iterations.", ConvergenceWarning)
C:\Users\Tanmay\Anaconda3\lib\site-packages\sklearn\svm\base.py:922: Converge
nceWarning: Liblinear failed to converge, increase the number of iterations.
    "the number of iterations.", ConvergenceWarning)
C:\Users\Tanmay\Anaconda3\lib\site-packages\sklearn\svm\base.py:922: Converge
nceWarning: Liblinear failed to converge, increase the number of iterations.
    "the number of iterations.", ConvergenceWarning)
C:\Users\Tanmay\Anaconda3\lib\site-packages\sklearn\svm\base.py:922: Converge
nceWarning: Liblinear failed to converge, increase the number of iterations.
    "the number of iterations.", ConvergenceWarning)
```

## Supervised Learning - Classification

1. KNN Classification
2. Logistic Regression
3. Linear SVM
4. Kernalized SVM
5. Decision Tree



```

In [63]: audit = pd.read_csv('audit_risk.csv')
        trial = pd.read_csv('trial.csv')
        audit['Money_Value'].fillna(0,inplace=True)
        trial['Money_Value'].fillna(0,inplace=True)
        # Sorting the data frames
        audit = audit.sort_values(by=['LOCATION_ID', 'TOTAL'])
        trial = trial.sort_values(by=['LOCATION_ID', 'TOTAL'])

        #Concatinate data frames
        data = pd.concat([audit,trial[['Loss', 'History_score', 'LOSS_SCORE', 'MONEY_Marks']]],axis=1)

        #Removing LOCATION_ID
        data.drop('LOCATION_ID',axis=1,inplace=True)

        data['TOTAL'].max()
        # 1268.91
        data[data['TOTAL'] == 1268.91]
        # index = 241
        data.drop(241,inplace=True)

```

Repeating the same Data pre-processing, exploratory analysis, and feature selection.

Only difference is for classification we consider Risk variable.

So we drop the Audit\_Risk and include Risk variable for analysis.

```

In [64]: data = data[['Sector_score', 'PARA_A', 'Score_A', 'PARA_B', 'Score_B', 'numbers',
        'Money_Value', 'Score_MV', 'District_Loss', 'LOSS_SCORE', 'History_score', 'Risk'
        ]]

```

```

In [65]: from sklearn.model_selection import train_test_split
        X = data.drop('Risk',axis=1)
        y = data['Risk']

        X_train_org, X_test_org, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=0)

        from sklearn.preprocessing import MinMaxScaler
        scale = MinMaxScaler()
        X_train = scale.fit_transform(X_train_org)
        X_test = scale.transform(X_test_org)

```

```

C:\Users\Tanmay\Anaconda3\lib\site-packages\sklearn\preprocessing\data.py:32
3: DataConversionWarning: Data with input dtype int64, float64 were all converted to float64 by MinMaxScaler.
    return self.partial_fit(X, y)

```


## 1. KNN Classification

```
In [66]: from sklearn.neighbors import KNeighborsClassifier
train_score = []
test_score = []

n = range(1,15)
for i in n:
    knn = KNeighborsClassifier(n_neighbors=i)
    knn.fit(X_train,y_train)
    train_score.append(knn.score(X_train,y_train))
    test_score.append(knn.score(X_test,y_test))

plt.plot(n,train_score,'b',label='Train score')
plt.plot(n,test_score,'r',label = 'Test score')
plt.legend()
```

Out[66]: <matplotlib.legend.Legend at 0x1e46da092b0>



```
In [67]: knn = KNeighborsClassifier(5)
knn.fit(X_train, y_train)
print('Train score: {:.4f}'.format(knn.score(X_train, y_train)))
print('Test score: {:.4f}'.format(knn.score(X_test, y_test)))
```

Train score: 0.9742  
Test score: 0.9613

```
In [68]: from sklearn.metrics import classification_report, confusion_matrix
predictions = knn.predict(X_test)
print(confusion_matrix(y_test, predictions))
print(classification_report(y_test, predictions))
```

```
[[87  1]
 [ 5 62]]
```

	precision	recall	f1-score	support
0	0.95	0.99	0.97	88
1	0.98	0.93	0.95	67
micro avg	0.96	0.96	0.96	155
macro avg	0.96	0.96	0.96	155
weighted avg	0.96	0.96	0.96	155

```
In [69]: #Cross validation
knn = KNeighborsClassifier(n_neighbors=5)
scores = cross_val_score(knn, X_train, y_train, cv=5, scoring='accuracy')
scores1 = cross_val_score(knn, X_test, y_test, cv=5, scoring='accuracy')

print('Avg train score: {:.4f}'.format(scores.mean()))
print('Avg train score: {:.4f}'.format(scores1.mean()))
```

Avg train score: 0.9629

Avg train score: 0.9348

```
In [70]: #Grid Search
from sklearn.model_selection import GridSearchCV
k_range = list(range(1, 21))
param_grid = dict(n_neighbors=k_range)
grid = GridSearchCV(knn, param_grid, cv=10, scoring='accuracy')
grid.fit(X_train, y_train)
print('Best Score: {:.4f}'.format(grid.best_score_))
print('Best Paramater: {}'.format(grid.best_params_))
```

Best Score: 0.9710

Best Paramater: {'n\_neighbors': 1}

```
In [71]: cvres = grid.cv_results_
for mean_score, params in zip(cvres["mean_test_score"], cvres["params"]):
    print(np.sqrt(mean_score), params)
```


```
0.9853769542340047 {'n_neighbors': 1}
0.9796312670361429 {'n_neighbors': 2}
0.981276324898574 {'n_neighbors': 3}
0.981276324898574 {'n_neighbors': 4}
0.9804541409880654 {'n_neighbors': 5}
0.9788077013024735 {'n_neighbors': 6}
0.981276324898574 {'n_neighbors': 7}
0.9771584874918804 {'n_neighbors': 8}
0.9796312670361429 {'n_neighbors': 9}
0.9779834420393966 {'n_neighbors': 10}
0.9788077013024735 {'n_neighbors': 11}
0.9763328358974787 {'n_neighbors': 12}
0.9771584874918804 {'n_neighbors': 13}
0.9755064854862865 {'n_neighbors': 14}
0.9746794344808963 {'n_neighbors': 15}
0.9738516810963533 {'n_neighbors': 16}
0.9746794344808963 {'n_neighbors': 17}
0.9730232235401101 {'n_neighbors': 18}
0.9730232235401101 {'n_neighbors': 19}
0.9713641887040999 {'n_neighbors': 20}
```

```
In [72]: import mglearn
x_b = X_train[0:774,[2,5]]
y_b = y_train[0:774]

knn = KNeighborsClassifier(7)
knn.fit(x_b, y_b)

mglearn.plots.plot_2d_separator(knn, x_b, fill=True, eps=0.5, alpha=.4)
mglearn.discrete_scatter(x_b[:, 0], x_b[:, 1], y_b)
```

```
Out[72]: [<matplotlib.lines.Line2D at 0x1e4718b6da0>,
<matplotlib.lines.Line2D at 0x1e471605198>]
```



## 2. Logistic Regression

```
In [73]: from sklearn.linear_model import LogisticRegression

c_range = [0.001, 0.01, 0.1, 1, 10, 100, 1000]
train_score_l1 = []
train_score_l2 = []
test_score_l1 = []
test_score_l2 = []

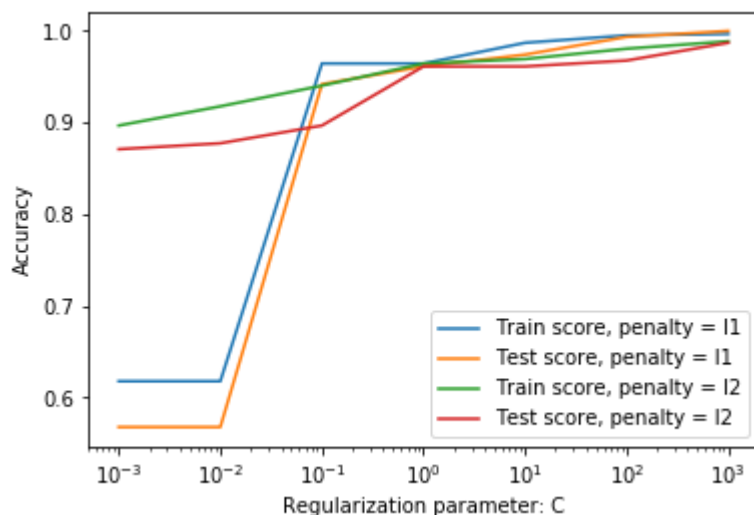
for c in c_range:
    log_l1 = LogisticRegression(penalty = 'l1', C = c)
    log_l2 = LogisticRegression(penalty = 'l2', C = c)
    log_l1.fit(X_train, y_train)
    log_l2.fit(X_train, y_train)
    train_score_l1.append(log_l1.score(X_train, y_train))
    train_score_l2.append(log_l2.score(X_train, y_train))
    test_score_l1.append(log_l1.score(X_test, y_test))
    test_score_l2.append(log_l2.score(X_test, y_test))

import matplotlib.pyplot as plt
%matplotlib inline

plt.plot(c_range, train_score_l1, label = 'Train score, penalty = l1')
plt.plot(c_range, test_score_l1, label = 'Test score, penalty = l1')
plt.plot(c_range, train_score_l2, label = 'Train score, penalty = l2')
plt.plot(c_range, test_score_l2, label = 'Test score, penalty = l2')
plt.legend()
plt.xlabel('Regularization parameter: C')
plt.ylabel('Accuracy')
plt.xscale('log')
```

```
C:\Users\Tanmay\Anaconda3\lib\site-packages\sklearn\linear_model\logistic.py:
433: FutureWarning: Default solver will be changed to 'lbfgs' in 0.22. Specif
y a solver to silence this warning.
FutureWarning)
C:\Users\Tanmay\Anaconda3\lib\site-packages\sklearn\linear_model\logistic.py:
433: FutureWarning: Default solver will be changed to 'lbfgs' in 0.22. Specif
y a solver to silence this warning.
FutureWarning)
C:\Users\Tanmay\Anaconda3\lib\site-packages\sklearn\linear_model\logistic.py:
433: FutureWarning: Default solver will be changed to 'lbfgs' in 0.22. Specif
y a solver to silence this warning.
FutureWarning)
C:\Users\Tanmay\Anaconda3\lib\site-packages\sklearn\linear_model\logistic.py:
433: FutureWarning: Default solver will be changed to 'lbfgs' in 0.22. Specif
y a solver to silence this warning.
FutureWarning)
C:\Users\Tanmay\Anaconda3\lib\site-packages\sklearn\linear_model\logistic.py:
433: FutureWarning: Default solver will be changed to 'lbfgs' in 0.22. Specif
y a solver to silence this warning.
FutureWarning)
C:\Users\Tanmay\Anaconda3\lib\site-packages\sklearn\linear_model\logistic.py:
433: FutureWarning: Default solver will be changed to 'lbfgs' in 0.22. Specif
y a solver to silence this warning.
FutureWarning)
C:\Users\Tanmay\Anaconda3\lib\site-packages\sklearn\linear_model\logistic.py:
433: FutureWarning: Default solver will be changed to 'lbfgs' in 0.22. Specif
y a solver to silence this warning.
FutureWarning)
C:\Users\Tanmay\Anaconda3\lib\site-packages\sklearn\linear_model\logistic.py:
433: FutureWarning: Default solver will be changed to 'lbfgs' in 0.22. Specif
y a solver to silence this warning.
FutureWarning)
C:\Users\Tanmay\Anaconda3\lib\site-packages\sklearn\linear_model\logistic.py:
433: FutureWarning: Default solver will be changed to 'lbfgs' in 0.22. Specif
y a solver to silence this warning.
FutureWarning)
C:\Users\Tanmay\Anaconda3\lib\site-packages\sklearn\linear_model\logistic.py:
433: FutureWarning: Default solver will be changed to 'lbfgs' in 0.22. Specif
y a solver to silence this warning.
FutureWarning)
C:\Users\Tanmay\Anaconda3\lib\site-packages\sklearn\linear_model\logistic.py:
433: FutureWarning: Default solver will be changed to 'lbfgs' in 0.22. Specif
y a solver to silence this warning.
FutureWarning)
C:\Users\Tanmay\Anaconda3\lib\site-packages\sklearn\linear_model\logistic.py:
433: FutureWarning: Default solver will be changed to 'lbfgs' in 0.22. Specif
y a solver to silence this warning.
FutureWarning)
C:\Users\Tanmay\Anaconda3\lib\site-packages\sklearn\linear_model\logistic.py:
433: FutureWarning: Default solver will be changed to 'lbfgs' in 0.22. Specif
y a solver to silence this warning.
FutureWarning)
C:\Users\Tanmay\Anaconda3\lib\site-packages\sklearn\linear_model\logistic.py:
433: FutureWarning: Default solver will be changed to 'lbfgs' in 0.22. Specif
y a solver to silence this warning.
FutureWarning)
C:\Users\Tanmay\Anaconda3\lib\site-packages\sklearn\svm\base.py:922: ConvergenceWarning: Liblinear failed to converge, increase the number of iterations.
"the number of iterations.", ConvergenceWarning)
C:\Users\Tanmay\Anaconda3\lib\site-packages\sklearn\linear_model\logistic.py:
433: FutureWarning: Default solver will be changed to 'lbfgs' in 0.22. Specif
```

y a solver to silence this warning.  
FutureWarning)



```
In [74]: log_l1 = LogisticRegression(penalty = 'l1', C = 0.1)
log_l2 = LogisticRegression(penalty = 'l2', C = 1)
log_l1.fit(X_train,y_train)
log_l2.fit(X_train,y_train)
print("Penalty: l1")
print("Train_score:{:.4f}".format(log_l1.score(X_train, y_train)))
print("Test_score:{:.4f}".format(log_l1.score(X_test, y_test)))
print("Penalty: l2")
print("Train_score:{:.4f}".format(log_l2.score(X_train, y_train)))
print("Test_score:{:.4f}".format(log_l2.score(X_test, y_test)))
```

```
Penalty: l1
Train_score:0.9645
Test_score:0.9419
Penalty: l2
Train_score:0.9645
Test_score:0.9613
```

```
C:\Users\tanmay\Anaconda3\lib\site-packages\sklearn\linear_model\logistic.py:
433: FutureWarning: Default solver will be changed to 'lbfgs' in 0.22. Specif
y a solver to silence this warning.
FutureWarning)
C:\Users\tanmay\Anaconda3\lib\site-packages\sklearn\linear_model\logistic.py:
433: FutureWarning: Default solver will be changed to 'lbfgs' in 0.22. Specif
y a solver to silence this warning.
FutureWarning)
```

```
In [75]: #Confusion Matrix
from sklearn.metrics import classification_report, confusion_matrix
pred = log_l2.predict(X_test)
print(confusion_matrix(y_test, pred))
print(classification_report(y_test, pred))
```

```
[[87  1]
 [ 5 62]]
```

	precision	recall	f1-score	support
0	0.95	0.99	0.97	88
1	0.98	0.93	0.95	67
micro avg	0.96	0.96	0.96	155
macro avg	0.96	0.96	0.96	155
weighted avg	0.96	0.96	0.96	155



```
In [76]: #Cross validation
log_l2 = LogisticRegression(penalty = 'l2', C = 1)
scores = cross_val_score(log_l2, X_train, y_train, cv=5, scoring='accuracy')
scores1= cross_val_score(log_l2, X_test, y_test, cv=5, scoring='accuracy')

print('Avg train score: {:.4f}'.format(scores.mean()))
print('Avg train score: {:.4f}'.format(scores1.mean()))
```

Avg train score: 0.9612

Avg train score: 0.9477

C:\Users\Tanmay\Anaconda3\lib\site-packages\sklearn\linear\_model\logistic.py:  
433: FutureWarning: Default solver will be changed to 'lbfgs' in 0.22. Specify a solver to silence this warning.

FutureWarning)

C:\Users\Tanmay\Anaconda3\lib\site-packages\sklearn\linear\_model\logistic.py:  
433: FutureWarning: Default solver will be changed to 'lbfgs' in 0.22. Specify a solver to silence this warning.

FutureWarning)

C:\Users\Tanmay\Anaconda3\lib\site-packages\sklearn\linear\_model\logistic.py:  
433: FutureWarning: Default solver will be changed to 'lbfgs' in 0.22. Specify a solver to silence this warning.

FutureWarning)

C:\Users\Tanmay\Anaconda3\lib\site-packages\sklearn\linear\_model\logistic.py:  
433: FutureWarning: Default solver will be changed to 'lbfgs' in 0.22. Specify a solver to silence this warning.

FutureWarning)

C:\Users\Tanmay\Anaconda3\lib\site-packages\sklearn\linear\_model\logistic.py:  
433: FutureWarning: Default solver will be changed to 'lbfgs' in 0.22. Specify a solver to silence this warning.

FutureWarning)

C:\Users\Tanmay\Anaconda3\lib\site-packages\sklearn\linear\_model\logistic.py:  
433: FutureWarning: Default solver will be changed to 'lbfgs' in 0.22. Specify a solver to silence this warning.

FutureWarning)

C:\Users\Tanmay\Anaconda3\lib\site-packages\sklearn\linear\_model\logistic.py:  
433: FutureWarning: Default solver will be changed to 'lbfgs' in 0.22. Specify a solver to silence this warning.

FutureWarning)

C:\Users\Tanmay\Anaconda3\lib\site-packages\sklearn\linear\_model\logistic.py:  
433: FutureWarning: Default solver will be changed to 'lbfgs' in 0.22. Specify a solver to silence this warning.

FutureWarning)

C:\Users\Tanmay\Anaconda3\lib\site-packages\sklearn\linear\_model\logistic.py:  
433: FutureWarning: Default solver will be changed to 'lbfgs' in 0.22. Specify a solver to silence this warning.

FutureWarning)

C:\Users\Tanmay\Anaconda3\lib\site-packages\sklearn\linear\_model\logistic.py:  
433: FutureWarning: Default solver will be changed to 'lbfgs' in 0.22. Specify a solver to silence this warning.

FutureWarning)


```
In [77]: import mglearn
x_b = X_train[50:100, [1,3]]
y_b = y_train[50:100]

lreg = LogisticRegression()
lreg.fit(x_b, y_b)

mglearn.plots.plot_2d_separator(lreg, x_b, fill=True, eps=0.5, alpha=.4)
mglearn.discrete_scatter(x_b[:, 0], x_b[:, 1], y_b)
```

C:\Users\tanmay\Anaconda3\lib\site-packages\sklearn\linear\_model\logistic.py:  
433: FutureWarning: Default solver will be changed to 'lbfgs' in 0.22. Specify a solver to silence this warning.  
FutureWarning)

```
Out[77]: [<matplotlib.lines.Line2D at 0x1e40012c5f8>,  
<matplotlib.lines.Line2D at 0x1e4001f4c50>]
```



### 3. Linear SVM

```
In [78]: from sklearn.svm import LinearSVC
linear_svm = LinearSVC()
linear_svm.fit(X_train,y_train)
print('Train score: {:.4f}'.format(linear_svm.score(X_train,y_train)))
print('Test score: {:.4f}'.format(linear_svm.score(X_test,y_test)))
```

Train score: 0.9710  
Test score: 0.9613

```
In [79]: #Confusion Matrix
from sklearn.metrics import classification_report, confusion_matrix
pred = linear_svm.predict(X_test)
print(confusion_matrix(y_test, pred))
print(classification_report(y_test, pred))
```

```
[[87  1]
 [ 5 62]]
```

	precision	recall	f1-score	support
0	0.95	0.99	0.97	88
1	0.98	0.93	0.95	67
micro avg	0.96	0.96	0.96	155
macro avg	0.96	0.96	0.96	155
weighted avg	0.96	0.96	0.96	155

```
In [80]: #Cross validation
linear_svm = LinearSVC()
scores = cross_val_score(linear_svm, X_train, y_train, cv=5, scoring='accuracy')
scores1 = cross_val_score(linear_svm, X_test, y_test, cv=5, scoring='accuracy')

print('Avg train score: {:.4f}'.format(scores.mean()))
print('Avg train score: {:.4f}'.format(scores1.mean()))
```

```
Avg train score: 0.9629
Avg train score: 0.9540
```

```
In [81]: mglearn.discrete_scatter(X_train[:, 0], X_train[:, 1], y_train)
line = np.linspace(-5, 5)
for coef, intercept, color in zip(linear_svm.coef_, linear_svm.intercept_,
                                mglearn.cm3.colors):
    plt.plot(line, -(line * coef[0] + intercept) / coef[1], c=color)
plt.ylim(-2, 2)
plt.xlim(0, 1.5)
plt.xlabel("Feature 0")
plt.ylabel("Feature 1")
plt.legend(['Class 0', 'Class 1', 'Line class 0', 'Line class 1'], loc=(1.01,
0.3))
```

-----

**AttributeError**

Traceback (most recent call last)

<ipython-input-81-28f5a8b24143> in <module>

1 mglearn.discrete\_scatter(X\_train[:, 0], X\_train[:, 1], y\_train)


2 line = np.linspace(-5, 5)

----> 3 for coef, intercept, color in zip(linear\_svm.coef\_, linear\_svm.intercept\_,

4 mglearn.cm3.colors):

5 plt.plot(line, -(line \* coef[0] + intercept) / coef[1], c=color)

**AttributeError:** 'LinearSVC' object has no attribute 'coef\_'



#### 4. Kernalized SVM

##### *RBF kernal*

```
In [83]: #kernel = 'rbf'
from sklearn.svm import SVC
C1 = [0.01,0.1,1,10]
gamma1 = [0.01,0.1,1,10]

for i in C1:
    for j in gamma1:
        svc = SVC(C=i,kernel='rbf',gamma=j)
        svc.fit(X_train,y_train)
        print('C:{},gamma:{}'.format(i,j))
        print('Train score: {:.4f},Test score: {:.4f}'.format(svc.score(X_train,y_train),svc.score(X_test,y_test)))
```

```
C:0.01,gamma:0.01
Train score: 0.6177,Test score: 0.5677
C:0.01,gamma:0.1
Train score: 0.6177,Test score: 0.5677
C:0.01,gamma:1
Train score: 0.9000,Test score: 0.8387
C:0.01,gamma:10
Train score: 0.6177,Test score: 0.5677
C:0.1,gamma:0.01
Train score: 0.6177,Test score: 0.5677
C:0.1,gamma:0.1
Train score: 0.9194,Test score: 0.8774
C:0.1,gamma:1
Train score: 0.9661,Test score: 0.9484
C:0.1,gamma:10
Train score: 0.9339,Test score: 0.9355
C:1,gamma:0.01
Train score: 0.9161,Test score: 0.8774
C:1,gamma:0.1
Train score: 0.9710,Test score: 0.9613
C:1,gamma:1
Train score: 0.9694,Test score: 0.9484
C:1,gamma:10
Train score: 0.9774,Test score: 0.9548
C:10,gamma:0.01
Train score: 0.9710,Test score: 0.9613
C:10,gamma:0.1
Train score: 0.9726,Test score: 0.9548
C:10,gamma:1
Train score: 0.9774,Test score: 0.9677
C:10,gamma:10
Train score: 0.9855,Test score: 0.9677
```

```
In [84]: #Cross validation
svc = SVC(C=0.1, kernel='rbf', gamma=1)
scores = cross_val_score(svc, X_train, y_train, cv=5, scoring='accuracy')
scores1 = cross_val_score(svc, X_test, y_test, cv=5, scoring='accuracy')

print('Avg train score: {:.4f}'.format(scores.mean()))
print('Avg train score: {:.4f}'.format(scores1.mean()))
```

Avg train score: 0.9629

Avg train score: 0.9544

Best Parameters: C = 0.1, gamma=1

### Linear Kernal

```
In [85]: #kernal = Linear
C1 = [0.01, 0.1, 1, 10]

for i in C1:
    svc = SVC(C=i, kernel='linear')
    svc.fit(X_train, y_train)
    print('C:{}'.format(i))
    print('Train score: {:.4f}, Test score: {:.4f}'.format(svc.score(X_train, y_train), svc.score(X_test, y_test)))
```

C:0.01

Train score: 0.8758, Test score: 0.8452

C:0.1

Train score: 0.9581, Test score: 0.9419

C:1

Train score: 0.9726, Test score: 0.9677

C:10

Train score: 0.9758, Test score: 0.9677

```
In [86]: #Cross validation
svc = SVC(C=10, kernel='linear')
scores = cross_val_score(svc, X_train, y_train, cv=5, scoring='accuracy')
scores1 = cross_val_score(svc, X_test, y_test, cv=5, scoring='accuracy')

print('Avg train score: {:.4f}'.format(scores.mean()))
print('Avg train score: {:.4f}'.format(scores1.mean()))
```

Avg train score: 0.9662

Avg train score: 0.9544

Best Paramter: C=1

### polynomial kernal

```
In [87]: #kernel = poly
C1 = [0.01,0.1,1,10]
gamma1 = [0.01,0.1,1,10]

for i in C1:
    for j in gamma1:
        svc = SVC(C=i,kernel='poly',gamma=j)
        svc.fit(X_train,y_train)
        print('C:{},gamma:{}'.format(i,j))
        print('Train score: {:.4f},Test score: {:.4f}'.format(svc.score(X_train,y_train),svc.score(X_test,y_test)))
```

```
C:0.01,gamma:0.01
Train score: 0.6177,Test score: 0.5677
C:0.01,gamma:0.1
Train score: 0.6177,Test score: 0.5677
C:0.01,gamma:1
Train score: 0.8806,Test score: 0.8194
C:0.01,gamma:10
Train score: 0.9790,Test score: 0.9548
C:0.1,gamma:0.01
Train score: 0.6177,Test score: 0.5677
C:0.1,gamma:0.1
Train score: 0.6177,Test score: 0.5677
C:0.1,gamma:1
Train score: 0.9548,Test score: 0.9226
C:0.1,gamma:10
Train score: 0.9903,Test score: 0.9806
C:1,gamma:0.01
Train score: 0.6177,Test score: 0.5677
C:1,gamma:0.1
Train score: 0.7855,Test score: 0.7613
C:1,gamma:1
Train score: 0.9677,Test score: 0.9548
C:1,gamma:10
Train score: 0.9952,Test score: 1.0000
C:10,gamma:0.01
Train score: 0.6177,Test score: 0.5677
C:10,gamma:0.1
Train score: 0.8806,Test score: 0.8194
C:10,gamma:1
Train score: 0.9790,Test score: 0.9548
C:10,gamma:10
Train score: 0.9952,Test score: 1.0000
```

```
In [88]: #Cross validation
svc = SVC(C=1,kernel='poly',gamma=10)
scores = cross_val_score(svc, X_train, y_train, cv=5, scoring='accuracy')
scores1= cross_val_score(svc, X_test, y_test, cv=5, scoring='accuracy')

print('Avg train score: {:.4f}'.format(scores.mean()))
print('Avg train score: {:.4f}'.format(scores1.mean()))
```

```
Avg train score: 0.9791
Avg train score: 0.9552
```

## 5. Decision Tree

```
In [89]: from sklearn.tree import DecisionTreeClassifier
dtree = DecisionTreeClassifier(max_depth=3,random_state=0)
dtree.fit(X_train, y_train)

print("Accuracy on training set: {:.3f}".format(dtree.score(X_train, y_train)))
print("Accuracy on test set: {:.3f}".format(dtree.score(X_test, y_test)))
```

Accuracy on training set: 0.956  
Accuracy on test set: 0.961

```
In [90]: #Cross validation
from sklearn.model_selection import cross_val_score
dtree = DecisionTreeClassifier(random_state=0)
scores = cross_val_score(dtree, X_train, y_train, cv=5, scoring='accuracy')
scores1= cross_val_score(dtree, X_test, y_test, cv=5, scoring='accuracy')

print('Avg train score: {:.4f}'.format(scores.mean()))
print('Avg test score: {:.4f}'.format(scores1.mean()))
```

Avg train score: 0.9742  
Avg test score: 0.9735


```
In [91]: #Cross validation
dtree = DecisionTreeClassifier(max_depth=3,random_state=0)
scores = cross_val_score(dtree, X_train, y_train, cv=5, scoring='accuracy')
scores1= cross_val_score(dtree, X_test, y_test, cv=5, scoring='accuracy')

print('Avg train score: {:.4f}'.format(scores.mean()))
print('Avg test score: {:.4f}'.format(scores1.mean()))
```

Avg train score: 0.9371  
Avg test score: 0.9544



```
In [92]: from sklearn import tree
dtree.fit(X_train, y_train)
def plot_feature_importances_risk(model):
    n_features = X.shape[1]
    plt.barh(range(n_features), model.feature_importances_, align='center')
    plt.yticks(np.arange(n_features), list(X.columns))
    plt.xlabel("Feature importance")
    plt.ylabel("Features")
    plt.ylim(-1, n_features)
plt.figure(figsize=(10,10))
plot_feature_importances_risk(dtree)
```



```
In [93]: import os
os.environ["PATH"] += os.pathsep + 'C:/Program Files (x86)/Graphviz2.38/bin/'
+ 'C:/Users/jmoha/AppData/Local/conda/conda/envs/fluffy/Lib/site-packages/graphviz' + 'C:/Users/jmoha/AppData/Local/conda/conda/envs/fluffy/Lib/site-packages/PIL'

# Create DOT data
from sklearn.tree import export_graphviz
dot_data = export_graphviz(dtrees, out_file=None, filled=True, rounded=True, special_characters=True)

# Draw graph
import pydotplus
graph = pydotplus.graph_from_dot_data(dot_data)

# Show graph
from PIL import *
import graphviz
from IPython.display import Image
Image(graph.create_png())
```

```
-----
ModuleNotFoundError                                Traceback (most recent call last)
<ipython-input-93-64e399046e3a> in <module>
      7
      8 # Draw graph
----> 9 import pydotplus
     10 graph = pydotplus.graph_from_dot_data(dot_data)
     11

ModuleNotFoundError: No module named 'pydotplus'
```

```
In [94]: #Confusion Matrix
from sklearn.metrics import classification_report, confusion_matrix
pred = dtrees.predict(X_test)
print(confusion_matrix(y_test, pred))
print(classification_report(y_test, pred))
```

```
[[88  0]
 [ 6 61]]
```

	precision	recall	f1-score	support
0	0.94	1.00	0.97	88
1	1.00	0.91	0.95	67
micro avg	0.96	0.96	0.96	155
macro avg	0.97	0.96	0.96	155
weighted avg	0.96	0.96	0.96	155

## Overview of Regression scores

All the below scores are cross validation scores

```
In [95]: Regression = {'Model':['Linear Regression','KNN Regression','Ridge Regression',
, 'Lasso Regression','SVM Regression'], 'Avg.Train Score':[0.7347,0.6872,0.6067,
0.7377,0.6581], 'Avg_Test_Score':[0.6588,0.4047,0.227,0.6695,0.4419]}
Regression_score = pd.DataFrame(Regression)
Regression_score
```

Out[95]:

	Model	Avg.Train Score	Avg_Test_Score
0	Linear Regression	0.7347	0.6588
1	KNN Regression	0.6872	0.4047
2	Ridge Regression	0.6067	0.2270
3	Lasso Regression	0.7377	0.6695
4	SVM Regression	0.6581	0.4419

## Overview of Classification scores

All the below scores are cross validation scores

```
In [96]: Classification = {'Model':['KNN classification','Logistic Regrerssion','Linear
SVM','SVC - rbf','SVC - linear','SVC - poly','Decision Tree'], 'Avg.Train Score':
[0.9629,0.9612,0.9629,0.9629,0.9662,0.9791,0.9726], 'Avg_Test_Score':[0.9348
,0.9477,0.9613,0.9544,0.9544,0.9552,0.9735]}
Classification_score = pd.DataFrame(Classification)
Classification_score
```

Out[96]:

	Model	Avg.Train Score	Avg_Test_Score
0	KNN classification	0.9629	0.9348
1	Logistic Regrerssion	0.9612	0.9477
2	Linear SVM	0.9629	0.9613
3	SVC - rbf	0.9629	0.9544
4	SVC - linear	0.9662	0.9544
5	SVC - poly	0.9791	0.9552
6	Decision Tree	0.9726	0.9735

## Result

**The Best Regression model: Lasso Regression**

Parameters: alpha = 0.001

## The Best Classification model: Decision Tree

*The important features in analysis : 'PARA\_A', 'PARA\_B', 'Score\_MV', 'District\_Loss', 'LOSS\_SCORE'*

In [ ]:

*Let's run the above models with important features only*

```
In [97]: p = data[['PARA_A', 'PARA_B', 'Score_MV', 'District_Loss', 'LOSS_SCORE']]
q = data['Risk']

from sklearn.model_selection import train_test_split

p_train_org, p_test_org, q_train, q_test = train_test_split(p, q, test_size=0.2, random_state=0)

from sklearn.preprocessing import MinMaxScaler
scale = MinMaxScaler()
p_train = scale.fit_transform(p_train_org)
p_test = scale.transform(p_test_org)
```

```
C:\Users\Tanmay\Anaconda3\lib\site-packages\sklearn\preprocessing\data.py:32
3: DataConversionWarning: Data with input dtype int64, float64 were all converted to float64 by MinMaxScaler.
    return self.partial_fit(X, y)
```

```


In [98]: # KNN Classifier
from sklearn.neighbors import KNeighborsClassifier
train_score = []
test_score = []

n = range(1,15)
for i in n:
    knn = KNeighborsClassifier(n_neighbors=i)
    knn.fit(p_train,q_train)
    train_score.append(knn.score(p_train,q_train))
    test_score.append(knn.score(p_test,q_test))

plt.plot(n,train_score,'b',label='Train score')
plt.plot(n,test_score,'r',label = 'Test score')
plt.legend()

```

Out[98]: <matplotlib.legend.Legend at 0x1e4001184e0>



```

In [ ]: knn = KNeighborsClassifier(3)
knn.fit(p_train, q_train)
print('Train score: {:.4f}'.format(knn.score(p_train,q_train)))
print('Test score: {:.4f}'.format(knn.score(p_test,q_test)))
predictions = knn.predict(p_test)
print(confusion_matrix(q_test,predictions))
print(classification_report(q_test,predictions))

```

```
In [100]: #Logistic Regression
from sklearn.linear_model import LogisticRegression


c_range = [0.001, 0.01, 0.1, 1, 10, 100, 1000]
train_score_l1 = []
train_score_l2 = []
test_score_l1 = []
test_score_l2 = []

for c in c_range:
    log_l1 = LogisticRegression(penalty = 'l1', C = c)
    log_l2 = LogisticRegression(penalty = 'l2', C = c)
    log_l1.fit(p_train, q_train)
    log_l2.fit(p_train, q_train)
    train_score_l1.append(log_l1.score(p_train, q_train))
    train_score_l2.append(log_l2.score(p_train, q_train))
    test_score_l1.append(log_l1.score(p_test, q_test))
    test_score_l2.append(log_l2.score(p_test, q_test))

import matplotlib.pyplot as plt
%matplotlib inline

plt.plot(c_range, train_score_l1, label = 'Train score, penalty = l1')
plt.plot(c_range, test_score_l1, label = 'Test score, penalty = l1')
plt.plot(c_range, train_score_l2, label = 'Train score, penalty = l2')
plt.plot(c_range, test_score_l2, label = 'Test score, penalty = l2')
plt.legend()
plt.xlabel('Regularization parameter: C')
plt.ylabel('Accuracy')
plt.xscale('log')
```

[illegible]





```
In [101]: log_l1 = LogisticRegression(penalty = 'l1', C = 0.1)
log_l2 = LogisticRegression(penalty = 'l2', C = 0.1)
log_l1.fit(p_train,q_train)
log_l2.fit(p_train,q_train)
print("Penalty: l1")
print('Train score: {:.4f}'.format(log_l1.score(p_train,q_train)))
print('Test score: {:.4f}'.format(log_l1.score(p_test,q_test)))
predictions = log_l1.predict(p_test)
print(confusion_matrix(q_test,predictions))
print(classification_report(q_test,predictions))
print("Penalty: l2")
print('Train score: {:.4f}'.format(log_l2.score(p_train,q_train)))
print('Test score: {:.4f}'.format(log_l2.score(p_test,q_test)))
predictions = log_l2.predict(p_test)
print(confusion_matrix(q_test,predictions))
print(classification_report(q_test,predictions))
```

Penalty: l1

Train score: 0.9113

Test score: 0.9161

[[88 0]

[13 54]]

	precision	recall	f1-score	support
0	0.87	1.00	0.93	88
1	1.00	0.81	0.89	67
micro avg	0.92	0.92	0.92	155
macro avg	0.94	0.90	0.91	155
weighted avg	0.93	0.92	0.91	155

Penalty: l2

Train score: 0.9113

Test score: 0.9161

[[88 0]

[13 54]]

	precision	recall	f1-score	support
0	0.87	1.00	0.93	88
1	1.00	0.81	0.89	67
micro avg	0.92	0.92	0.92	155
macro avg	0.94	0.90	0.91	155
weighted avg	0.93	0.92	0.91	155

C:\Users\tanmay\Anaconda3\lib\site-packages\sklearn\linear\_model\logistic.py:  
433: FutureWarning: Default solver will be changed to 'lbfgs' in 0.22. Specify a solver to silence this warning.

FutureWarning)

C:\Users\tanmay\Anaconda3\lib\site-packages\sklearn\linear\_model\logistic.py:  
433: FutureWarning: Default solver will be changed to 'lbfgs' in 0.22. Specify a solver to silence this warning.

FutureWarning)

```
In [102]: # RBF SVM
#kernel = 'rbf'
C1 = [0.01,0.1,1,10]
gamma1 = [0.01,0.1,1,10]

for i in C1:
    for j in gamma1:
        svc = SVC(C=i,kernel='rbf',gamma=j)
        svc.fit(p_train,q_train)
        print('C:{},gamma:{}'.format(i,j))
        print('Train score: {:.4f},Test score: {:.4f}'.format(svc.score(p_train,q_train),svc.score(p_test,q_test)))
```

```
C:0.01,gamma:0.01
Train score: 0.6177,Test score: 0.5677
C:0.01,gamma:0.1
Train score: 0.6177,Test score: 0.5677
C:0.01,gamma:1
Train score: 0.8823,Test score: 0.8774
C:0.01,gamma:10
Train score: 0.9032,Test score: 0.8903
C:0.1,gamma:0.01
Train score: 0.6177,Test score: 0.5677
C:0.1,gamma:0.1
Train score: 0.8290,Test score: 0.8065
C:0.1,gamma:1
Train score: 0.9145,Test score: 0.9161
C:0.1,gamma:10
Train score: 0.9274,Test score: 0.9097
C:1,gamma:0.01
Train score: 0.8290,Test score: 0.8065
C:1,gamma:0.1
Train score: 0.9210,Test score: 0.9161
C:1,gamma:1
Train score: 0.9274,Test score: 0.9097
C:1,gamma:10
Train score: 0.9548,Test score: 0.9355
C:10,gamma:0.01
Train score: 0.9242,Test score: 0.9226
C:10,gamma:0.1
Train score: 0.9306,Test score: 0.9097
C:10,gamma:1
Train score: 0.9581,Test score: 0.9548
C:10,gamma:10
Train score: 0.9742,Test score: 0.9742
```

```
In [103]: #Decision Tree
from sklearn.tree import DecisionTreeClassifier
dtree = DecisionTreeClassifier(max_depth=3,random_state=0)
dtree.fit(p_train, q_train)

print("Accuracy on training set: {:.3f}".format(dtree.score(p_train, q_train)))
print("Accuracy on test set: {:.3f}".format(dtree.score(p_test, q_test)))

predictions = dtree.predict(p_test)
print(confusion_matrix(q_test,predictions))
print(classification_report(q_test,predictions))
```

Accuracy on training set: 0.952

Accuracy on test set: 0.948

[[87 1]

[ 7 60]]

	precision	recall	f1-score	support
0	0.93	0.99	0.96	88
1	0.98	0.90	0.94	67
micro avg	0.95	0.95	0.95	155
macro avg	0.95	0.94	0.95	155
weighted avg	0.95	0.95	0.95	155

In [ ]: